

2327422-code

September 3, 2024

0.0.1 1. Load the dataset in Jupyter Notebook using pandas

```
[1]: import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset
df = pd.read_csv("Dataset (1).csv")

# Display the first few rows of the dataset
df.head()
```

```
[1]: SLNO  Candidate Ref DOJ Extended Duration to accept offer Notice period \
0      1      2110407      Yes      14.00      30.00
1      2      2112635      No      18.00      30.00
2      3      2112838      No      3.00      45.00
3      4      2115021      No      26.00      30.00
4      5      2115125      Yes      1.00      120.00
```

```
Offered band Pecent hike expected in CTC Percent hike offered in CTC \
0      E2      -20.79      13.16
1      E2      50.00      320.00
2      E2      42.84      42.84
3      E2      42.84      42.84
4      E2      42.59      42.59
```

```
Percent difference CTC Joining Bonus Candidate relocate actual Gender \
0      42.86      No      No      Female
1      180.00      No      No      Male
2      0.00      No      No      Male
3      0.00      No      No      Male
4      0.00      No      Yes      Male
```

```
Candidate Source Rex in Yrs LOB Location Age Status
0      Agency      7.0      ERS      Noida      34.00      Joined
1      Employee Referral      8.0      INFRA      Chennai      34.00      Joined
2      Agency      4.0      INFRA      Noida      27.00      Joined
3      Employee Referral      4.0      INFRA      Noida      34.00      Joined
4      Employee Referral      6.0      INFRA      Noida      34.00      Joined
```

```
[2]: import numpy as np

# Replace empty strings or spaces with NaN
df['Duration to accept offer'] = df['Duration to accept offer'].replace(' ', np.
    ↪nan)
df['Notice period'] = df['Notice period'].replace(' ', np.nan)
df['Age'] = df['Age'].replace(' ', np.nan)

# Now try converting to float
df['Duration to accept offer'] = df['Duration to accept offer'].astype(float)
df['Notice period'] = df['Notice period'].astype(float)
df['Age'] = df['Age'].astype(float)
```

```
[3]: df.dtypes
```

```
[3]: SLNO                                int64
Candidate Ref                          int64
DOJ Extended                          object
Duration to accept offer               float64
Notice period                         float64
Offered band                          object
Pecent hike expected in CTC           object
Percent hike offered in CTC           object
Percent difference CTC                object
Joining Bonus                        object
Candidate relocate actual             object
Gender                               object
Candidate Source                     object
Rex in Yrs                           float64
LOB                                  object
Location                             object
Age                                  float64
Status                               object
dtype: object
```

```
[4]: # Convert columns to numeric, setting errors='coerce' to turn non-numeric_
    ↪values into NaN
df['Pecent hike expected in CTC'] = pd.to_numeric(df['Pecent hike expected in_
    ↪CTC'], errors='coerce')
df['Percent hike offered in CTC'] = pd.to_numeric(df['Percent hike offered in_
    ↪CTC'], errors='coerce')
df['Percent difference CTC'] = pd.to_numeric(df['Percent difference CTC'],_
    ↪errors='coerce')

# Fill NaN values with the mean of the column
```

```
df['Pecent hike expected in CTC'] = df['Pecent hike expected in CTC'].
    ↪fillna(df['Pecent hike expected in CTC'].mean())
df['Percent hike offered in CTC'] = df['Percent hike offered in CTC'].
    ↪fillna(df['Percent hike offered in CTC'].mean())
df['Percent difference CTC'] = df['Percent difference CTC'].fillna(df['Percent_
    ↪difference CTC'].mean())
```

```
[5]: # Summarize the number of null values in each column
null_summary = df.isnull().sum()

# Display the summary
print(null_summary)
```

```
SLNO          0
Candidate Ref  0
DOJ Extended  0
Duration to accept offer  4
Notice period  3
Offered band   0
Pecent hike expected in CTC  0
Percent hike offered in CTC  0
Percent difference CTC  0
Joining Bonus  0
Candidate relocate actual  0
Gender         0
Candidate Source  0
Rex in Yrs     0
LOB           0
Location       0
Age           2
Status        0
dtype: int64
```

```
[6]: df.describe()
```

```
[6]:
```

	SLNO	Candidate Ref	Duration to accept offer	Notice period \
count	8995.000000	8.995000e+03	8991.000000	8992.000000
mean	5970.984325	2.843647e+06	21.434768	39.291593
std	3373.963454	4.863448e+05	25.809203	22.223567
min	1.000000	2.109586e+06	0.000000	0.000000
25%	3207.500000	2.386476e+06	3.000000	30.000000
50%	5976.000000	2.807482e+06	10.000000	30.000000
75%	8739.000000	3.300060e+06	33.000000	60.000000
max	12333.000000	3.836076e+06	224.000000	120.000000

	Pecent hike expected in CTC	Percent hike offered in CTC \
count	8995.000000	8995.000000

mean	43.868558	40.666311
std	29.778215	36.052575
min	-68.830000	-60.530000
25%	27.270000	22.200000
50%	40.000000	36.080000
75%	53.850000	50.000000
max	359.770000	471.430000

	Percent difference CTC	Rex in Yrs	Age
count	8995.000000	8995.000000	8993.000000
mean	-1.600444	4.239022	29.913933
std	19.493676	2.547571	4.097990
min	-67.270000	0.000000	20.000000
25%	-8.330000	3.000000	27.000000
50%	0.000000	4.000000	29.000000
75%	0.000000	6.000000	34.000000
max	300.000000	24.000000	60.000000

0.0.2 2. Build a correlation matrix between all the numeric features in the dataset. Report the features, which are correlated at a cut-off of 0.70. What actions will you take on the features, which are highly correlated?


```
[7]: # Get the list of numerical feature labels
numerical_features = df.select_dtypes(include=[np.number]).columns

# Display the numerical feature labels
print(numerical_features)
```

```
Index(['SLNO', 'Candidate Ref', 'Duration to accept offer', 'Notice period',
      'Pecent hike expected in CTC', 'Percent hike offered in CTC',
      'Percent difference CTC', 'Rex in Yrs', 'Age'],
      dtype='object')
```

```
[8]: import seaborn as sns
import matplotlib.pyplot as plt

correlation_matrix = df[['Duration to accept offer', 'Notice period',
                        'Pecent hike expected in CTC', 'Percent hike offered_
                        ↳in CTC',
                        'Percent difference CTC', 'Rex in Yrs', 'Age']].corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f',
↳linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()
```



Highest Correlated Values are Percent Hike Offered in CTC and Percent Hike Expected in CTC with correlation of 0.67.

Actions on Highly Correlated Features: Feature Elimination:

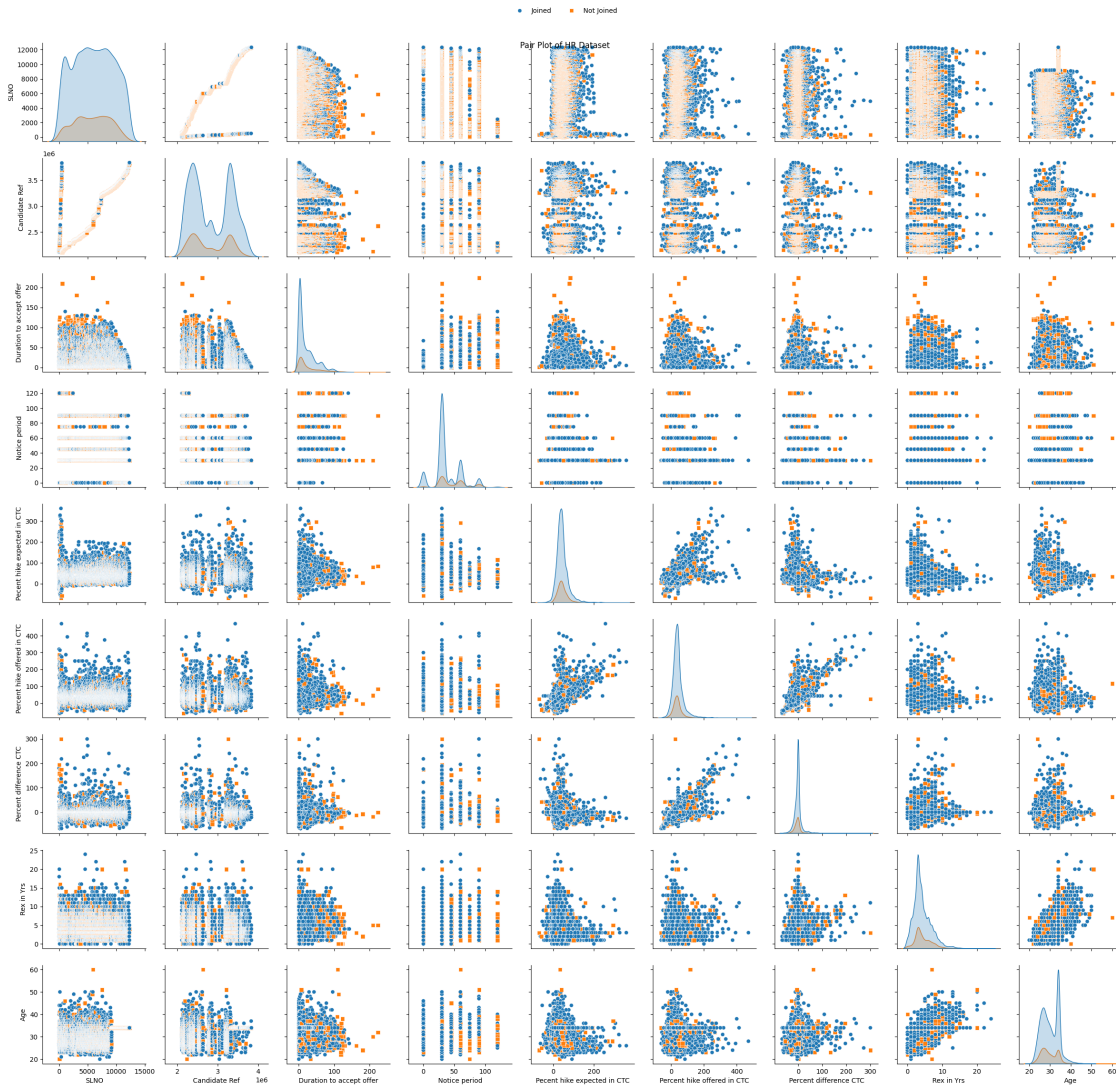
If two features are highly correlated, we might consider removing one of them to reduce redundancy. Keeping both can lead to multicollinearity in some models, which can affect model performance. Dimensionality Reduction: Instead of removing, We can also apply dimensionality reduction techniques like Principal Component Analysis (PCA) to combine correlated features into a single component. Regularization Techniques: If we want to keep all features, consider using regularization techniques like Lasso or Ridge regression, which can handle multicollinearity by penalizing large coefficients.

```
[9]: import seaborn as sns
ax = sns.pairplot(data=df, hue='Status', markers=["o", "s"])
plt.suptitle("Pair Plot of HR Dataset")
sns.move_legend(
```

```

ax, "lower center",
bbox_to_anchor=(.5, 1), ncol=3, title=None, frameon=False)
plt.tight_layout()
plt.show()

```

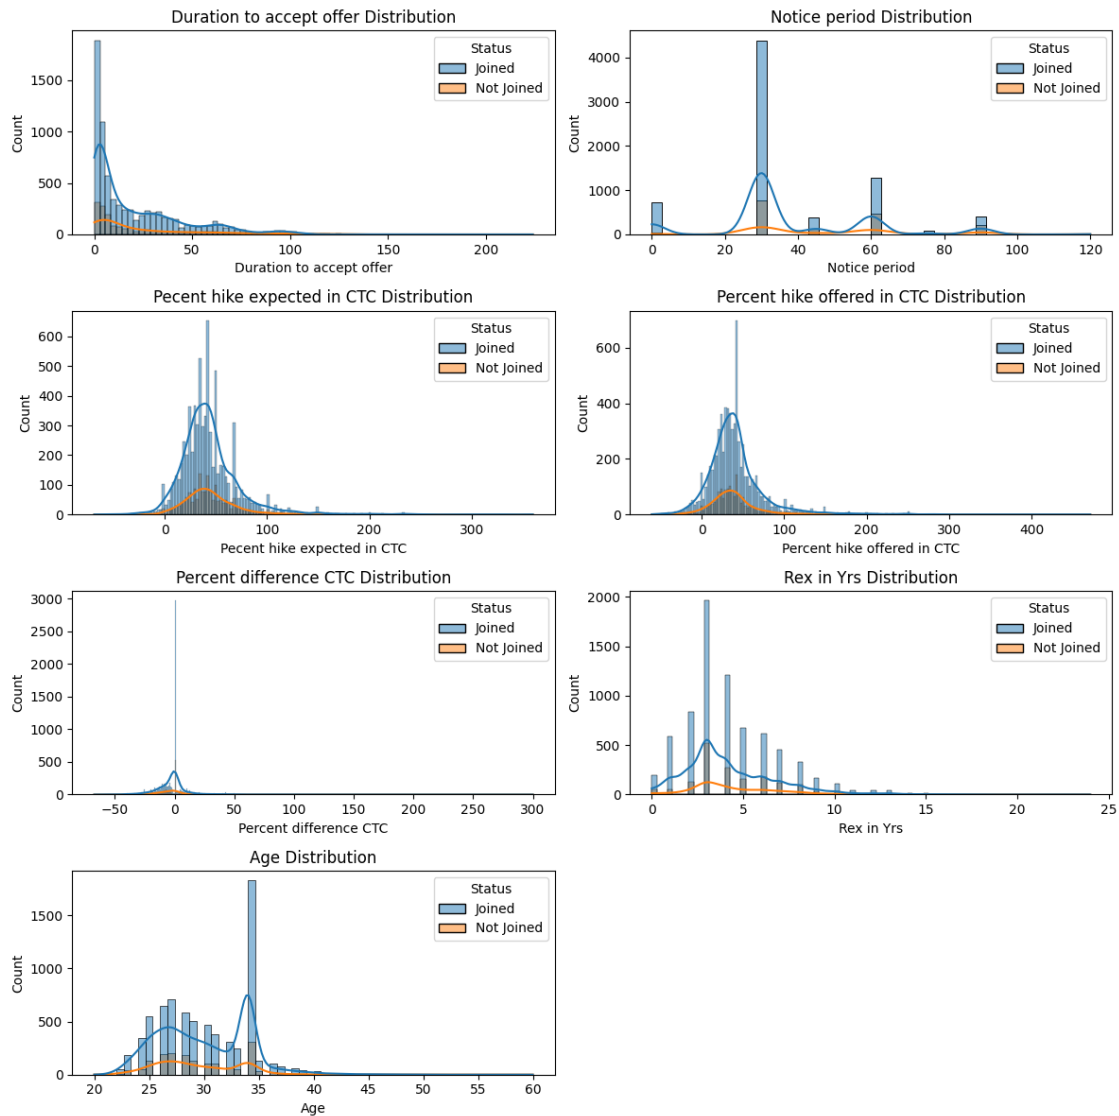


```

[10]: cols = ['Duration to accept offer', 'Notice period', 'Percent hike_
↳ expected in CTC', 'Percent hike offered in CTC', 'Percent_
↳ difference CTC', 'Rex in Yrs', 'Age']
# Visualize the distribution of each feature using histograms.
plt.figure(figsize=(12, 12))
for i, feature in enumerate(cols):
    plt.subplot(4, 2, i + 1)
    sns.histplot(data=df, x=feature, hue='Status', kde=True)
    plt.title(f'{feature} Distribution')

```

```
plt.tight_layout()
plt.show()
```



```
[11]: df['Status'] = df['Status'].replace({'Joined': 1, 'Not Joined': 0})
```

```
C:\Users\haris\AppData\Local\Temp\ipykernel_13616\1064645211.py:1:
FutureWarning: Downcasting behavior in `replace` is deprecated and will be
removed in a future version. To retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
df['Status'] = df['Status'].replace({'Joined': 1, 'Not Joined': 0})
```

0.0.3 3. Build a new feature named LOB_Hike_Offered using LOB and percentage hike offered. Include this as a part of the data frame created in step 1. What assumption are you trying to test with such variables?

```
[12]: df['LOB_Hike_Offered'] = df['LOB'] + '_' + df['Percent hike offered in CTC'].
      ↪astype(str)
```

```
[13]: df.head()
```

```
[13]:
```

	SLNO	Candidate Ref	DOJ	Extended	Duration to accept offer	Notice period	\
0	1	2110407		Yes	14.0	30.0	
1	2	2112635		No	18.0	30.0	
2	3	2112838		No	3.0	45.0	
3	4	2115021		No	26.0	30.0	
4	5	2115125		Yes	1.0	120.0	

	Offered band	Percent hike expected in CTC	Percent hike offered in CTC	\
0	E2	-20.79	13.16	
1	E2	50.00	320.00	
2	E2	42.84	42.84	
3	E2	42.84	42.84	
4	E2	42.59	42.59	

	Percent difference CTC	Joining Bonus	Candidate relocate	actual	Gender	\
0	42.86	No	No	No	Female	
1	180.00	No	No	No	Male	
2	0.00	No	No	No	Male	
3	0.00	No	No	No	Male	
4	0.00	No	Yes	Yes	Male	

	Candidate Source	Rex in Yrs	LOB	Location	Age	Status	\
0	Agency	7.0	ERS	Noida	34.0	1	
1	Employee Referral	8.0	INFRA	Chennai	34.0	1	
2	Agency	4.0	INFRA	Noida	27.0	1	
3	Employee Referral	4.0	INFRA	Noida	34.0	1	
4	Employee Referral	6.0	INFRA	Noida	34.0	1	

	LOB_Hike_Offered
0	ERS_13.16
1	INFRA_320.0
2	INFRA_42.84
3	INFRA_42.84
4	INFRA_42.59

```
[14]: df.value_counts(['Status'])
```



```
[14]: Status
      1      7313
      0      1682
      Name: count, dtype: int64
```

The new feature LOB_Hike_Offered tests the assumption that the hike offered for a particular line of business (LOB) has a significant impact on the outcome variable. By creating this feature, you can examine whether the combination of LOB and the percentage hike offered influences the target variable, such as whether a candidate joins the company or not.

0.0.4 4. Create a new data frame with the numeric features and categorical features as dummy variable coded features. Which features will you include for model building and why?

```
[15]: labels=df.columns
      print(labels)

Index(['SLNO', 'Candidate Ref', 'DOJ Extended', 'Duration to accept offer',
      'Notice period', 'Offered band', 'Pecent hike expected in CTC',
      'Percent hike offered in CTC', 'Percent difference CTC',
      'Joining Bonus', 'Candidate relocate actual', 'Gender',
      'Candidate Source', 'Rex in Yrs', 'LOB', 'Location', 'Age', 'Status',
      'LOB_Hike_Offered'],
      dtype='object')
```

Features That will be included for model building are 'DOJ Extended', 'Duration to accept offer', 'Notice period', 'Offered band', 'Pecent hike expected in CTC', 'Percent hike offered in CTC', 'Percent difference CTC', 'Joining Bonus', 'Candidate relocate actual', 'Gender', 'Candidate Source', 'Rex in Yrs', 'LOB', 'Location', 'Age'. All the Variables have a impact in deciding the candidates preference.

```
[16]: df.drop(columns=['SLNO', 'Candidate Ref', 'LOB_Hike_Offered'], inplace=True)
```

```
[17]: # Convert categorical variables into dummy/indicator variables
      df_dummies = pd.get_dummies(df, drop_first=True)

      # Display the first few rows of the updated DataFrame
      df_dummies.head()
```

```
[17]:   Duration to accept offer  Notice period  Pecent hike expected in CTC \
0                14.0             30.0                -20.79
1                18.0             30.0                 50.00
2                 3.0             45.0                 42.84
3                26.0             30.0                 42.84
4                 1.0            120.0                 42.59

      Percent hike offered in CTC  Percent difference CTC  Rex in Yrs  Age \
```

0		13.16		42.86	7.0	34.0
1		320.00		180.00	8.0	34.0
2		42.84		0.00	4.0	27.0
3		42.84		0.00	4.0	34.0
4		42.59		0.00	6.0	34.0

	Status	DOJ	Extended_Yes	Offered	band_EO	...	Location_Bangalore	\
0	1		True		False	...	False	
1	1		False		False	...	False	
2	1		False		False	...	False	
3	1		False		False	...	False	
4	1		True		False	...	False	

	Location_Chennai	Location_Cochin	Location_Gurgaon	Location_Hyderabad	\
0	False	False	False	False	False
1	True	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False

	Location_Kolkata	Location_Mumbai	Location_Noida	Location_Others	\
0	False	False	True	False	
1	False	False	False	False	
2	False	False	True	False	
3	False	False	True	False	
4	False	False	True	False	

	Location_Pune
0	False
1	False
2	False
3	False
4	False

[5 rows x 42 columns]

```
[18]: df_dummies = df_dummies.dropna()
```

```
[19]: from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split

# Define the features and target variable
X = df_dummies.drop('Status', axis=1)
y = df_dummies['Status']

from imblearn.over_sampling import SMOTE
```

```
smote=SMOTE(sampling_strategy='minority')
X,y=smote.fit_resample(X,y)
y.value_counts()
```

```
[19]: Status
1    7305
0    7305
Name: count, dtype: int64
```

0.0.5 5. Split the data into training set and test set. Use 80% of data for model training and 20% for model testing.

```
[20]: # Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

```
[21]: print(df_dummies.columns)
```

```
Index(['Duration to accept offer', 'Notice period',
      'Percent hike expected in CTC', 'Percent hike offered in CTC',
      'Percent difference CTC', 'Rex in Yrs', 'Age', 'Status',
      'DOJ_Extended_Yes', 'Offered band_E0', 'Offered band_E1',
      'Offered band_E2', 'Offered band_E3', 'Joining Bonus_No',
      'Joining Bonus_Yes', 'Candidate relocate actual_No',
      'Candidate relocate actual_Yes', 'Gender_Female', 'Gender_Male',
      'Candidate Source_Agency', 'Candidate Source_Direct',
      'Candidate Source_Employee Referral', 'LOB_AXON', 'LOB_BFSI',
      'LOB_CSMP', 'LOB_EAS', 'LOB_ERS', 'LOB_ETS', 'LOB_Healthcare',
      'LOB_INFRA', 'LOB_MMS', 'Location_Ahmedabad', 'Location_Bangalore',
      'Location_Chennai', 'Location_Cochin', 'Location_Gurgaon',
      'Location_Hyderabad', 'Location_Kolkata', 'Location_Mumbai',
      'Location_Noida', 'Location_Others', 'Location_Pune'],
      dtype='object')
```

```
[22]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# Initialize the logistic regression model
model = LogisticRegression(random_state=42)

# Train the model using the balanced dataset
model.fit(X, y)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
```

```

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Optional: Get the coefficients of the model
coefficients = model.coef_[0]
features = X.columns

# Create a DataFrame of coefficients
coef_df = pd.DataFrame({
    'Feature': features,
    'Coefficient': coefficients
})

print("\nFeature Coefficients:")
print(coef_df)

```

Confusion Matrix:

```

[[1157  298]
 [ 289 1178]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.80	0.80	1455
1	0.80	0.80	0.80	1467
accuracy			0.80	2922
macro avg	0.80	0.80	0.80	2922
weighted avg	0.80	0.80	0.80	2922

Feature Coefficients:

	Feature	Coefficient
0	Duration to accept offer	0.004409
1	Notice period	-0.012674
2	Pecent hike expected in CTC	0.001686
3	Percent hike offered in CTC	0.005351
4	Percent difference CTC	-0.000090
5	Rex in Yrs	-0.211960
6	Age	0.334844
7	DOJ Extended_Yes	-0.514626
8	Offered band_E0	-0.684236
9	Offered band_E1	-0.757630
10	Offered band_E2	-1.724475

11	Offered band_E3	-0.516516
12	Joining Bonus_No	0.103092
13	Joining Bonus_Yes	-0.580903
14	Candidate relocate actual_No	-1.483326
15	Candidate relocate actual_Yes	2.036789
16	Gender_Female	-1.017498
17	Gender_Male	-1.121320
18	Candidate Source_Agency	-2.287927
19	Candidate Source_Direct	-2.147763
20	Candidate Source_Employee Referral	-0.890152
21	LOB_AXON	-1.226850
22	LOB_BFSI	-1.496234
23	LOB_CSMP	-0.946933
24	LOB_EAS	-0.985091
25	LOB_ERS	-1.245858
26	LOB_ETS	-0.721098
27	LOB_Healthcare	-0.183444
28	LOB_INFRA	-0.825032
29	LOB_MMS	0.049331
30	Location_Ahmedabad	-0.023380
31	Location_Bangalore	-1.756293
32	Location_Chennai	-1.996664
33	Location_Cochin	0.001259
34	Location_Gurgaon	-0.312187
35	Location_Hyderabad	-0.748390
36	Location_Kolkata	-0.483066
37	Location_Mumbai	-0.198432
38	Location_Noida	-1.544684
39	Location_Others	0.021498
40	Location_Pune	-0.094502

C:\Users\haris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

0.0.6 6. Build a model using Gender and Age as independent variable and Status as dependent variable.

- Are Gender and Age a significant feature in this model?

- What inferences can be drawn from this model?

```
[23]: from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split

# Define the features and target variable
X = df_dummies[['Gender_Female', 'Gender_Male', 'Age']]
y = df_dummies['Status']

# Split the data into 80% training and 20% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Apply SMOTE to balance the training dataset
smote = SMOTE(sampling_strategy='minority', random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

# Check the distribution of the target variable after SMOTE
y_train_smote.value_counts()
```

```
[23]: Status
1      5825
0      5825
Name: count, dtype: int64
```

Without SMOTE

```
[40]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# Initialize the logistic regression model
model = LogisticRegression(random_state=42)

# Train the model using the training data
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

```
Confusion Matrix:
[[ 46  449]
```

```
[ 28 2173]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.62	0.09	0.16	495
1	0.83	0.99	0.90	2201
accuracy			0.82	2696
macro avg	0.73	0.54	0.53	2696
weighted avg	0.79	0.82	0.77	2696

C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

With SMOTE

```
[24]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# Initialize the logistic regression model
model = LogisticRegression(random_state=42)

# Train the model using the SMOTE-balanced training data
model.fit(X_train_smote, y_train_smote)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

#
```

Confusion Matrix:

```
[[180 138]
 [719 761]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.20	0.57	0.30	318
1	0.85	0.51	0.64	1480
accuracy			0.52	1798
macro avg	0.52	0.54	0.47	1798
weighted avg	0.73	0.52	0.58	1798

The Variables Gender and Age are one of the significant features in the model. The Models that they have significantly higher Precision and recall Value on Joined they contribute well along with other Variables.

0.0.7 7. Build a model with statsmodel.api to predict the probability of Not Joining. How do you interpret the model outcome? Report the model performance on the test set.

```
[25]: import statsmodels.api as sm

X_array = np.asarray(X_train, dtype=float)
y_array = np.asarray(y_train, dtype=int) # Assuming your dependent variable is
↳ binary (0 or 1)

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_array)

logit_model=sm.Logit(y_array, X_array)
result=logit_model.fit()

print(result.summary())
```

Optimization terminated successfully.

Current function value: 0.484582

Iterations 6

Logit Regression Results

```
=====
Dep. Variable:                y    No. Observations:                7188
Model:                        Logit    Df Residuals:                7185
Method:                        MLE    Df Model:                    2
```



```

Date:                Tue, 03 Sep 2024    Pseudo R-squ.:          0.002248
Time:                23:10:36           Log-Likelihood:         -3483.2
converged:           True               LL-Null:                -3491.0
Covariance Type:     nonrobust          LLR p-value:            0.0003906
=====

```

	coef	std err	z	P> z	[0.025	0.975]
x1	0.7066	0.223	3.168	0.002	0.269	1.144
x2	0.5775	0.223	2.594	0.009	0.141	1.014
x3	0.0287	0.007	3.884	0.000	0.014	0.043

=====

```

[45]: from sklearn.linear_model import LogisticRegression
      from sklearn import metrics
      logreg = LogisticRegression()
      logreg.fit(X_train, y_train)

```

C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```

[45]: LogisticRegression()

```

```

[46]: y_pred = logreg.predict(X_test)

```

```

[47]: print('Accuracy of logistic regression classifier on test set : {:.2f}'.
      ↪format(logreg.score(X_test, y_test)))

```

Accuracy of logistic regression classifier on test set : 0.82

```

[48]: from sklearn import model_selection
      from sklearn.model_selection import cross_val_score
      kfold = model_selection.KFold(n_splits=10, shuffle=True, random_state=7)
      modelCV = LogisticRegression()
      scoring = 'accuracy'
      results = model_selection.cross_val_score(modelCV, X_train, y_train, cv=kfold,
      ↪scoring=scoring)
      print("10-fold cross validation average accuracy: {:.3f}" % (results.mean()))

```

```
C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\harris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\haris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

10-fold cross validation average accuracy: 0.811

```
C:\Users\haris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression


```
n_iter_i = _check_optimize_result(
```

```
[49]: from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[ 46 449]
 [ 28 2173]]
```

```
[50]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, logreg.predict_proba(X_test)[:,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
```

```
plt.savefig('Log_ROC')
plt.show()
```



Coming the Logit Model with statsmodel.api the accuracy is 0.82 and Results: Precision for class 0: ~0.0929 Recall for class 0: ~0.6216

with 10 cross validation accuracy is 0.811

The Model is Signifacant in identifying not joined candidates

- 1 # 8. Build a model with statsmodel.formula.api to predict the probability of Not Joining and report the model performance on the test set. What difference do you observe in the model built here and the one built in step 7.

```
[ ]: import pandas as pd
import statsmodels.formula.api as smf
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, roc_auc_score
```

```

# Load the dataset
df = pd.read_csv("Dataset (1).csv")

# Preprocess the Data
df['Not_Joined'] = df['Status'].apply(lambda x: 1 if x == 'Not Joined' else 0)
df = df.drop(columns=['SLNO', 'Candidate Ref', 'Status'])

# Sanitize column names (if necessary)
df.columns = df.columns.str.replace(' ', '_').str.replace(r'[^a-zA-Z0-9_]', '',
↪ regex=True)

# Split the data into training and testing sets
train_df, test_df = train_test_split(df, test_size=0.2, random_state=42)

# Build the Logistic Regression Model using the formula API
formula = 'Not_Joined ~ ' + ' + '.join([col for col in df.columns if col !=
↪ 'Not_Joined'])
model = smf.logit(formula=formula, data=train_df)
result = model.fit()

# Display the model summary
print(result.summary())

# Evaluate the Model on the Test Set
y_pred_prob = result.predict(test_df)
y_pred = (y_pred_prob >= 0.5).astype(int)
y_test = test_df['Not_Joined']

# Calculate performance metrics
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

# Display the performance metrics
print(f'Accuracy: {accuracy:.2f}')
print('Confusion Matrix:')
print(conf_matrix)

```

```

/usr/local/lib/python3.10/dist-
packages/statsmodels/discrete/discrete_model.py:2385: RuntimeWarning: overflow
encountered in exp
    return 1/(1+np.exp(-X))

```

True Positives and True Negatives: The model is good at predicting “Not Joining” with 1436 correct predictions. However, it struggles with correctly predicting “Joining,” with only 21 correct predictions. False Negatives: The model has a high number of false negatives (319), which indicates

that it often fails to identify candidates who will actually join. This is a significant issue if predicting “Joining” is crucial. False Positives: The model has a relatively low number of false positives (23), which means it rarely predicts “Joining” when the candidate won’t join

The model has a high accuracy of 81%, but this is mainly driven by its ability to predict “Not Joining” correctly. The model struggles significantly with predicting “Joining,” as indicated by the low true positives and high false negatives.

2 9. Build a model using sklearn package to predict the probability of Not Joining. What difference do you observe in this model compared to model built in step 7 and 8.

With SMOTE

```
[31]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    mean_squared_error

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.metrics import accuracy_score, classification_report

#LDA
lda_model = LinearDiscriminantAnalysis()
lda_model.fit(X_train_smote, y_train_smote)
# Make predictions
lda_predictions = lda_model.predict(X_test)
# Evaluate the model
print("LDA Accuracy:", accuracy_score(y_test, lda_predictions))
print("LDA Classification Report:\n", classification_report(y_test, \
    lda_predictions))

# Logistic Regression
lr_model = LogisticRegression()
lr_model.fit(X_train_smote, y_train_smote)
lr_predictions = lr_model.predict(X_test)
print("Logistic Regression Accuracy:", accuracy_score(y_test, lr_predictions))
print("Logistic Regression Classification Report:\n", \
    classification_report(y_test, lr_predictions))

# Multiple Linear Regression
mlr_model = LinearRegression()
mlr_model.fit(X_train_smote, y_train_smote)
mlr_predictions = mlr_model.predict(X_test)
```

```

print("Multiple Linear Regression Mean Squared Error:",
      ↪mean_squared_error(y_test, mlr_predictions))

# Random Forest
rf_model = RandomForestClassifier()
rf_model.fit(X_train_smote, y_train_smote)
rf_predictions = rf_model.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, rf_predictions))
print("Random Forest Classification Report:\n", classification_report(y_test,
      ↪rf_predictions))

# Decision Tree
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train_smote, y_train_smote)
dt_predictions = dt_model.predict(X_test)
print("Decision Tree Accuracy:", accuracy_score(y_test, dt_predictions))
print("Decision Tree Classification Report:\n", classification_report(y_test,
      ↪dt_predictions))

from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC

# K-Nearest Neighbors
knn_model = KNeighborsClassifier()
knn_model.fit(X_train_smote, y_train_smote)
knn_predictions = knn_model.predict(X_test)
print("K-Nearest Neighbors Accuracy:", accuracy_score(y_test, knn_predictions))
print("K-Nearest Neighbors Classification Report:\n",
      ↪classification_report(y_test, knn_predictions))

# Support Vector Machine
svm_model = SVC()
svm_model.fit(X_train_smote, y_train_smote)
svm_predictions = svm_model.predict(X_test)
print("Support Vector Machine Accuracy:", accuracy_score(y_test,
      ↪svm_predictions))
print("Support Vector Machine Classification Report:\n",
      ↪classification_report(y_test, svm_predictions))

```

LDA Accuracy: 0.5233592880978866

LDA Classification Report:

	precision	recall	f1-score	support
0	0.20	0.57	0.30	318
1	0.85	0.51	0.64	1480
accuracy			0.52	1798

macro avg	0.52	0.54	0.47	1798
weighted avg	0.73	0.52	0.58	1798

Logistic Regression Accuracy: 0.5233592880978866

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.20	0.57	0.30	318
1	0.85	0.51	0.64	1480
accuracy			0.52	1798
macro avg	0.52	0.54	0.47	1798
weighted avg	0.73	0.52	0.58	1798

Multiple Linear Regression Mean Squared Error: 0.24820502225071117

Random Forest Accuracy: 0.5278086763070078

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.20	0.54	0.29	318
1	0.84	0.53	0.65	1480
accuracy			0.53	1798
macro avg	0.52	0.53	0.47	1798
weighted avg	0.73	0.53	0.58	1798

Decision Tree Accuracy: 0.5216907675194661

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.20	0.56	0.29	318
1	0.85	0.51	0.64	1480
accuracy			0.52	1798
macro avg	0.52	0.54	0.47	1798
weighted avg	0.73	0.52	0.58	1798

K-Nearest Neighbors Accuracy: 0.7202447163515017

K-Nearest Neighbors Classification Report:

	precision	recall	f1-score	support
0	0.20	0.20	0.20	318
1	0.83	0.83	0.83	1480
accuracy			0.72	1798
macro avg	0.52	0.52	0.52	1798
weighted avg	0.72	0.72	0.72	1798

Support Vector Machine Accuracy: 0.4582869855394883

Support Vector Machine Classification Report:

	precision	recall	f1-score	support
0	0.20	0.69	0.31	318
1	0.86	0.41	0.55	1480
accuracy			0.46	1798
macro avg	0.53	0.55	0.43	1798
weighted avg	0.74	0.46	0.51	1798

Without SMOTE

```
[34]: X = df_dummies.drop('Status', axis=1)
      y = df_dummies['Status']

      # Split the data
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)

[35]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression, LinearRegression
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import accuracy_score, classification_report,
      ↪mean_squared_error

      #LDA
      lda_model = LinearDiscriminantAnalysis()
      lda_model.fit(X_train, y_train)
      # Make predictions
      lda_predictions = lda_model.predict(X_test)
      # Evaluate the model
      print("LDA Accuracy:", accuracy_score(y_test, lda_predictions))
      print("LDA Classification Report:\n", classification_report(y_test,
      ↪lda_predictions))

      # Logistic Regression
      lr_model = LogisticRegression()
      lr_model.fit(X_train, y_train)
      lr_predictions = lr_model.predict(X_test)
      print("Logistic Regression Accuracy:", accuracy_score(y_test, lr_predictions))
      print("Logistic Regression Classification Report:\n",
      ↪classification_report(y_test, lr_predictions))

      # Multiple Linear Regression
```

```

mlr_model = LinearRegression()
mlr_model.fit(X_train, y_train)
mlr_predictions = mlr_model.predict(X_test)
print("Multiple Linear Regression Mean Squared Error:",
      ↪mean_squared_error(y_test, mlr_predictions))

# Random Forest
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, rf_predictions))
print("Random Forest Classification Report:\n",
      ↪classification_report(y_test,
                             ↪rf_predictions))

# Decision Tree
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train, y_train)
dt_predictions = dt_model.predict(X_test)
print("Decision Tree Accuracy:", accuracy_score(y_test, dt_predictions))
print("Decision Tree Classification Report:\n",
      ↪classification_report(y_test,
                             ↪dt_predictions))

from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC

# K-Nearest Neighbors
knn_model = KNeighborsClassifier()
knn_model.fit(X_train, y_train)
knn_predictions = knn_model.predict(X_test)
print("K-Nearest Neighbors Accuracy:", accuracy_score(y_test, knn_predictions))
print("K-Nearest Neighbors Classification Report:\n",
      ↪classification_report(y_test, knn_predictions))

# Support Vector Machine
svm_model = SVC()
svm_model.fit(X_train, y_train)
svm_predictions = svm_model.predict(X_test)
print("Support Vector Machine Accuracy:", accuracy_score(y_test,
      ↪svm_predictions))
print("Support Vector Machine Classification Report:\n",
      ↪classification_report(y_test, svm_predictions))

```

LDA Accuracy: 0.8275862068965517

LDA Classification Report:

	precision	recall	f1-score	support
0	0.62	0.06	0.11	318

1	0.83	0.99	0.90	1480
accuracy			0.83	1798
macro avg	0.73	0.53	0.51	1798
weighted avg	0.79	0.83	0.76	1798

Logistic Regression Accuracy: 0.8275862068965517

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.58	0.09	0.16	318
1	0.84	0.99	0.90	1480
accuracy			0.83	1798
macro avg	0.71	0.54	0.53	1798
weighted avg	0.79	0.83	0.77	1798

C:\Users\haris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Multiple Linear Regression Mean Squared Error: 0.13044942533556403

Random Forest Accuracy: 0.8286985539488321

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.55	0.16	0.25	318
1	0.84	0.97	0.90	1480
accuracy			0.83	1798
macro avg	0.70	0.57	0.58	1798
weighted avg	0.79	0.83	0.79	1798

Decision Tree Accuracy: 0.7730812013348165

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.36	0.36	0.36	318
1	0.86	0.86	0.86	1480

accuracy			0.77	1798
macro avg	0.61	0.61	0.61	1798
weighted avg	0.77	0.77	0.77	1798

K-Nearest Neighbors Accuracy: 0.7914349276974416

K-Nearest Neighbors Classification Report:

	precision	recall	f1-score	support
0	0.32	0.16	0.21	318
1	0.84	0.93	0.88	1480

accuracy			0.79	1798
macro avg	0.58	0.54	0.55	1798
weighted avg	0.75	0.79	0.76	1798

Support Vector Machine Accuracy: 0.8236929922135706

Support Vector Machine Classification Report:

	precision	recall	f1-score	support
0	1.00	0.00	0.01	318
1	0.82	1.00	0.90	1480

accuracy			0.82	1798
macro avg	0.91	0.50	0.45	1798
weighted avg	0.85	0.82	0.74	1798

Naive Bayes

WITHOUT SMOTE

```
[36]: from sklearn.naive_bayes import GaussianNB
      from sklearn.metrics import accuracy_score, classification_report

      # Create and train the Gaussian Naive Bayes model
      gnb_model = GaussianNB()
      gnb_model.fit(X_train, y_train)

      # Make predictions
      gnb_predictions = gnb_model.predict(X_test)

      # Evaluate the model
      print("Gaussian Naive Bayes Accuracy:", accuracy_score(y_test, gnb_predictions))
      print("Gaussian Naive Bayes Classification Report:\n",
            ↪classification_report(y_test, gnb_predictions))
```

Gaussian Naive Bayes Accuracy: 0.3264738598442714

Gaussian Naive Bayes Classification Report:

	precision	recall	f1-score	support
0	0.21	0.99	0.34	318
1	0.99	0.18	0.31	1480
accuracy			0.33	1798
macro avg	0.60	0.59	0.33	1798
weighted avg	0.85	0.33	0.32	1798

With SMOTE

```
[37]: from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report
from imblearn.over_sampling import SMOTE

# Example preprocessing
# Encode categorical features if necessary (e.g., using LabelEncoder or
#   ↳ OneHotEncoder)
# Ensure X_train_smote and X_test have the same features after preprocessing

# Standardize features if needed
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Apply SMOTE to the training data
smote = SMOTE()
X_train_smote, y_train_smote = smote.fit_resample(X_train_scaled, y_train)

# Create and train the Gaussian Naive Bayes model
gnb_model = GaussianNB()
gnb_model.fit(X_train_smote, y_train_smote)

# Make predictions
gnb_predictions = gnb_model.predict(X_test_scaled)

# Evaluate the model
print("Gaussian Naive Bayes Accuracy:", accuracy_score(y_test, gnb_predictions))
print("Gaussian Naive Bayes Classification Report:\n",
#   ↳ classification_report(y_test, gnb_predictions))
```

Gaussian Naive Bayes Accuracy: 0.3264738598442714

Gaussian Naive Bayes Classification Report:

	precision	recall	f1-score	support
0	0.21	0.99	0.34	318

1	0.99	0.18	0.31	1480
accuracy			0.33	1798
macro avg	0.60	0.59	0.33	1798
weighted avg	0.85	0.33	0.32	1798

Interpretation of Models Without SMOTE: Accuracy: Most models show high accuracy, with Random Forest achieving the highest (0.8287). However, accuracy alone can be misleading in imbalanced datasets. Precision for “Not Joined” (Class 0): Precision is relatively higher in models like LDA (0.62) and SVM (1.00), indicating that when these models predict “Not Joined,” they are more likely to be correct. Recall for “Not Joined” (Class 0): Recall is generally low across models, with SVM having a recall of 0.00, indicating that it fails to capture any true “Not Joined” cases. Decision Tree shows a better balance with a recall of 0.36, meaning it captures 36% of the actual “Not Joined” cases.

*Interpretation of Models With SMOTE: Accuracy: Accuracy drops for most models, with K-Nearest Neighbors (KNN) showing the highest accuracy of 0.7202 after SMOTE. Precision for “Not Joined” (Class 0): Precision is uniform across all models (0.20), reflecting a more balanced approach towards predicting the “Not Joined” class but at the cost of accuracy. Recall for “Not Joined” (Class 0): Recall is higher across models compared to the results without SMOTE, with Support Vector Machine (SVM) achieving the highest recall (0.69). This suggests that SMOTE helped the models better capture the minority class, but the cost was reduced precision and accuracy.

Overall Interpretation: Without SMOTE: The models achieve higher accuracy but at the cost of recall, particularly for the minority class (Not Joined). The SVM model achieves perfect precision but fails to recall any minority cases, making it ineffective for balanced prediction. With SMOTE:

SMOTE helps in improving recall significantly, but at the cost of accuracy and precision. The models now identify more “Not Joined” cases but also produce more false positives. Conclusion: SMOTE is effective in improving the ability of the models to recognize the minority class (“Not Joined”), as evidenced by the improved recall scores. However, this comes with a trade-off in accuracy and precision. Depending on the specific goals of the model (e.g., whether it’s more important to correctly identify “Not Joined” candidates), using SMOTE might be a better option despite the lower accuracy. If precision is critical, the models without SMOTE may be preferable.

2.1 10. Fine-tune the cut-off value using cost of misclassification as a strategy. The cut-off should help classify maximum number of Not Joining cases correctly.

2.2 11.Fine-tune the cut-off value using youdens index as a strategy. The cut-off should help balance the classification of Joined and Not Joined cases.

```
[51]: from sklearn.model_selection import train_test_split

from sklearn.metrics import confusion_matrix
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```

# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42)

# Train a logistic regression model (replace with your model)
model = LogisticRegression()
model.fit(X_train, y_train)

# Get predicted probabilities
y_pred_sklearn_prob = model.predict_proba(X_test)[:, 1]

# Define a function to calculate cost of misclassification and Youden's Index
def calculate_cut_off(y_true, y_prob, cut_off):
    y_pred = (y_prob >= cut_off).astype(int)
    tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()
    cost_of_misclassification = fp + fn
    # Prevent division by zero
    youdens_index = tp / (tp + fn) - fp / (fp + tn) if (tp + fn) and (fp + tn)
↳else 0
    return cost_of_misclassification, youdens_index

# Iterate over different cut-off values
cut_off_values = [0.1 * i for i in range(1, 10)]
misclassification_costs = []
youdens_indices = []

for cut_off in cut_off_values:
    cost, youden = calculate_cut_off(y_test, y_pred_sklearn_prob, cut_off)
    misclassification_costs.append(cost)
    youdens_indices.append(youden)

# Find the optimal cut-offs
optimal_cut_off_cost = cut_off_values[misclassification_costs.
↳index(min(misclassification_costs))]
optimal_cut_off_youden = cut_off_values[youdens_indices.
↳index(max(youdens_indices))]

print(f'Optimal Cut-Off (Cost of Misclassification): {optimal_cut_off_cost}')
print(f'Optimal Cut-Off (Youden\'s Index): {optimal_cut_off_youden}')

# Plotting the results
plt.figure(figsize=(12, 6))
plt.plot(cut_off_values, misclassification_costs, label='Cost of
↳Misclassification', marker='o')
plt.plot(cut_off_values, youdens_indices, label='Youden\'s Index', marker='o')
plt.axvline(x=optimal_cut_off_cost, color='r', linestyle='--', label='Optimal
↳Cut-Off (Cost)')

```



```
plt.axvline(x=optimal_cut_off_youden, color='g', linestyle='--', label='Optimal_
↳Cut-Off (Youden\'s Index)')
plt.xlabel('Cut-Off Value')
plt.ylabel('Metric Value')
plt.title('Cut-Off Value vs Metrics')
plt.legend()
plt.grid(True)
plt.show()
```

C:\Users\haris\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

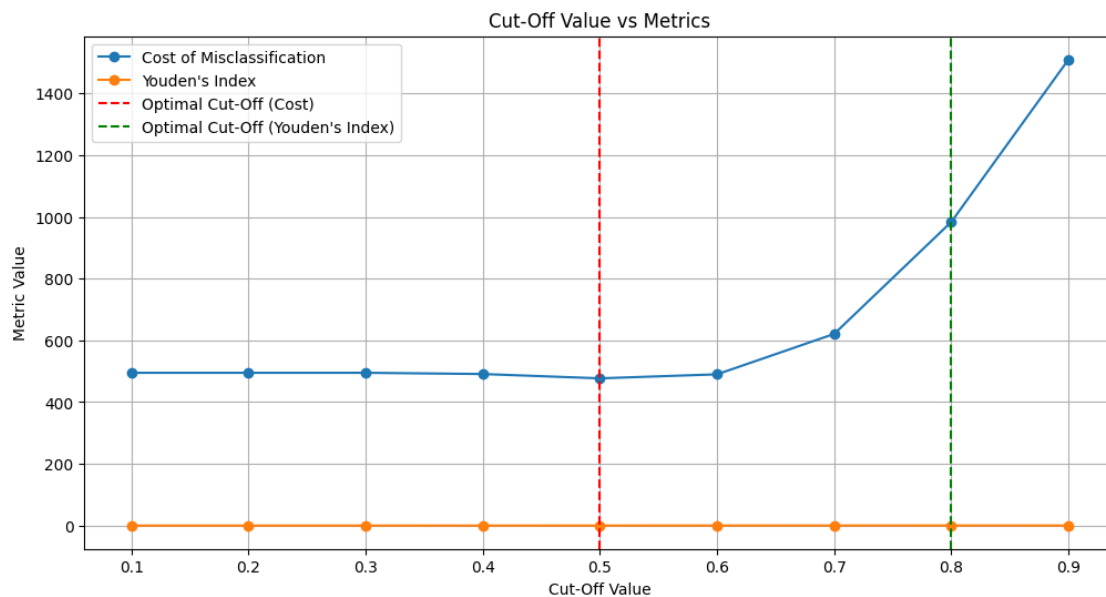
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Optimal Cut-Off (Cost of Misclassification): 0.5

Optimal Cut-Off (Youden's Index): 0.8



1. Optimal Cut-Off (Cost of Misclassification): 0.5 Interpretation: A cut-off of 0.5 is the default threshold used in many binary classification models. This means that if a model predicts a probability of 0.5 or higher, the observation is classified into the positive class (e.g., “Joined”).

This threshold is chosen to balance the cost of false positives and false negatives equally. However, this might not always be the best threshold, especially in cases where the cost of misclassifications is not equal (e.g., predicting “Not Joined” when someone actually “Joined” might have different consequences compared to the reverse).

2. Optimal Cut-Off (Youden’s Index): 0.8 Interpretation: Youden’s Index is a measure that seeks to maximize the difference between the true positive rate (sensitivity) and the false positive rate (1-specificity). It’s calculated as: $\text{Youden's Index} = \text{Sensitivity} + \text{Specificity} - 1$. A cut-off of 0.8 indicates that this threshold maximizes Youden’s Index, meaning it provides the best trade-off between sensitivity and specificity for this model. A higher threshold like 0.8 suggests that only predictions with a high probability (80% or more) are classified into the positive class. This would typically result in fewer false positives but might increase false negatives.

Implications: Cut-Off of 0.5 (Cost of Misclassification): This is often a good starting point, particularly when the costs of different types of errors (false positives vs. false negatives) are similar. However, it may not be the optimal choice if the cost of misclassification is significantly different for the two classes. Cut-Off of 0.8 (Youden’s Index):

This threshold is more conservative, classifying fewer cases as positive. It’s optimal for situations where you want to minimize false positives and are okay with potentially missing some true positives. It suggests that the model performs better (in terms of balancing sensitivity and specificity) when a higher threshold is used, which might be useful in situations where false positives are costly.

2.3 12. Apply the cut-off values obtained in step 10 and step 11 on the test set. What inference can be deduced from it?

```
[39]: from sklearn.metrics import classification_report, confusion_matrix

# Apply the optimal cut-offs
y_pred_cost = (y_pred_sklearn_prob >= optimal_cut_off_cost).astype(int)
y_pred_youden = (y_pred_sklearn_prob >= optimal_cut_off_youden).astype(int)

# Calculate confusion matrices
cm_cost = confusion_matrix(y_test, y_pred_cost)
cm_youden = confusion_matrix(y_test, y_pred_youden)

# Print confusion matrices
print(f'Confusion Matrix (Cost of Misclassification):\n{cm_cost}')
print(f'Confusion Matrix (Youden\'s Index):\n{cm_youden}')

# Print classification reports for detailed metrics
print(f'Classification Report (Cost of Misclassification):
↪\n{classification_report(y_test, y_pred_cost)}')
print(f'Classification Report (Youden\'s Index):
↪\n{classification_report(y_test, y_pred_youden)}')
```

Confusion Matrix (Cost of Misclassification):

```
[[ 46  449]
 [ 28 2173]]
```

Confusion Matrix (Youden's Index):

```
[[ 354  141]
 [ 843 1358]]
```

Classification Report (Cost of Misclassification):

	precision	recall	f1-score	support
0	0.62	0.09	0.16	495
1	0.83	0.99	0.90	2201
accuracy			0.82	2696
macro avg	0.73	0.54	0.53	2696
weighted avg	0.79	0.82	0.77	2696

Classification Report (Youden's Index):

	precision	recall	f1-score	support
0	0.30	0.72	0.42	495
1	0.91	0.62	0.73	2201
accuracy			0.64	2696
macro avg	0.60	0.67	0.58	2696
weighted avg	0.79	0.64	0.68	2696

Key Takeaways: Cost of Misclassification:

High Accuracy (0.82), but at the cost of low recall for the “Not Joined” class (0.09). This model is heavily skewed towards correctly predicting “Joined” but struggles significantly with identifying “Not Joined” cases. The model’s high accuracy is misleading as it fails to identify a significant number of “Not Joined” cases. Youden’s Index:

Lower Accuracy (0.64), but with a more balanced approach. This model improves recall for the “Not Joined” class (0.72) but sacrifices some recall for the “Joined” class (0.62). This model is better suited if you value a more balanced identification of both classes, even if it means sacrificing overall accuracy. Conclusion: Cost of Misclassification: This threshold is useful if the primary goal is to avoid missing out on predicting the “Joined” cases, but it will miss many “Not Joined” cases. Youden’s Index: This threshold provides a more balanced approach, capturing a higher proportion of “Not Joined” cases at the cost of reduced overall accuracy. This might be more suitable if both classes are important in the decision-making process.